

# dog\_app

April 30, 2020

## 1 Convolutional Neural Networks

### 1.1 Project: Write an Algorithm for a Dog Identification App

---

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '(IMPLEMENTATION)' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

**Note:** Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

**Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

## Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets:

**Note:** if you are using the Udacity workspace, you *DO NOT* need to re-download these - they can be found in the /data folder as noted in the cell below.

- Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location `/dog_images`.
- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location `/lfw`.

*Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.*

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays `human_files` and `dog_files`.

```
In [1]: import numpy as np
       from glob import glob

       # load filenames for human and dog images
       human_files = np.array(glob("/data/lfw/*/*"))
       dog_files = np.array(glob("/data/dog_images/*/*/*"))

       # print number of images in each dataset
       print('There are %d total human images.' % len(human_files))
       print('There are %d total dog images.' % len(dog_files))

There are 13233 total human images.
There are 8351 total dog images.
```

### ## Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the `haarcascades` directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [2]: import cv2
       import matplotlib.pyplot as plt
       %matplotlib inline

       # extract pre-trained face detector
       face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

       # load color (BGR) image
       img = cv2.imread(human_files[0])
       # convert BGR image to grayscale
       gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

       # find faces in image
       faces = face_cascade.detectMultiScale(gray)

       # print number of faces detected in the image
       print('Number of faces detected:', len(faces))
```

```

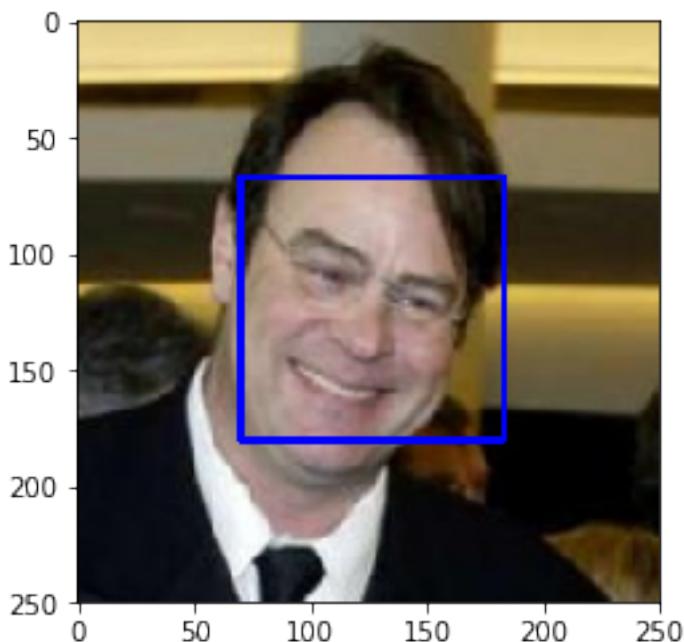
# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()

```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

### 1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns True if a human face is detected in an image and False otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [3]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    return len(faces) > 0
```

### 1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

**Question 1:** Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

**Answer:** - Percentage of human images with a detected human face is: 98 - Percentage of dog images with a detected human face is: 17

```
In [4]: from tqdm import tqdm

human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

#-#-# Do NOT modify the code above this line. #-#-#

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.
detected_human = 0
detected_dog = 0
for i in range(len(human_files_short)):
    if face_detector(human_files_short[i]):
        detected_human +=1
    if face_detector(dog_files_short[i]):
        detected_dog +=1

print("Percentage of human images with a detected human face is:", detected_human)
print("Percentage of dog images with a detected human face is:", detected_dog)
```

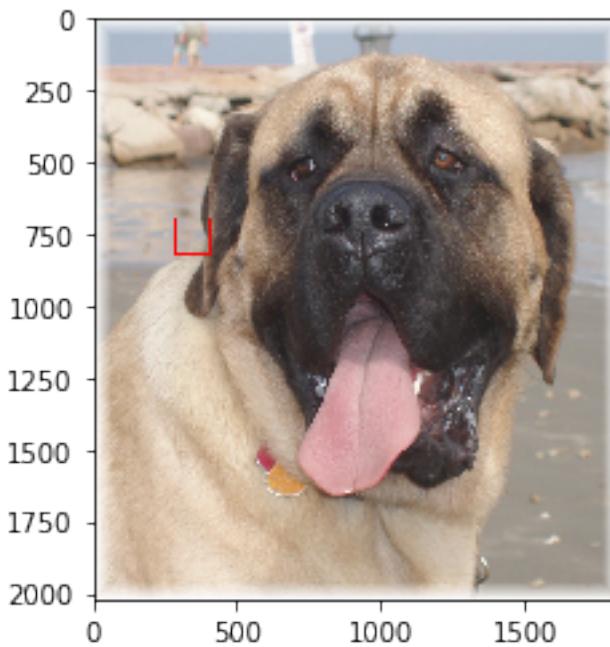
Percentage of human images with a detected human face is: 98

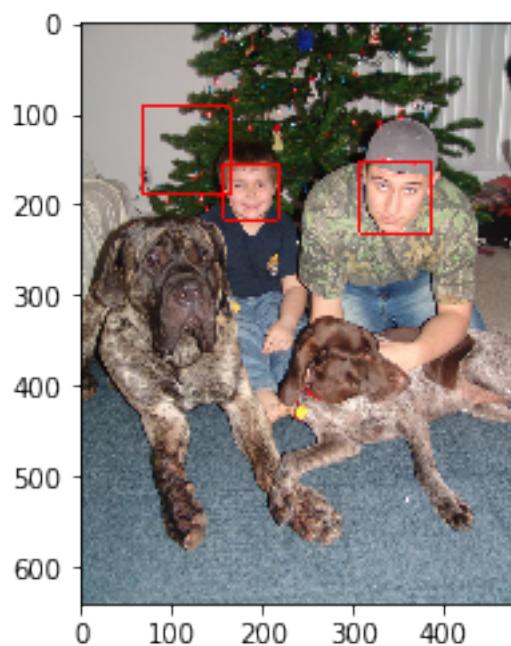
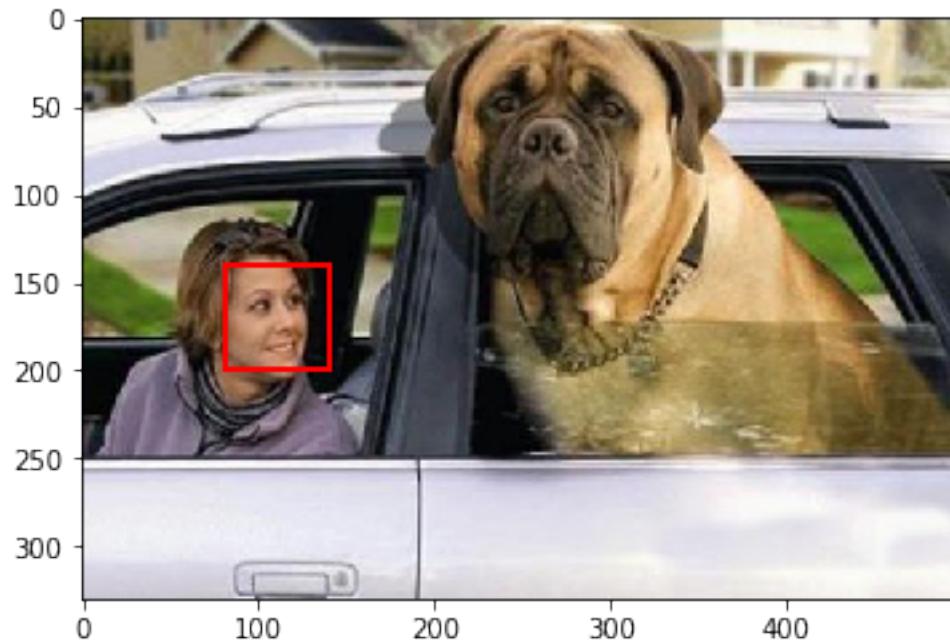
Percentage of dog images with a detected human face is: 17

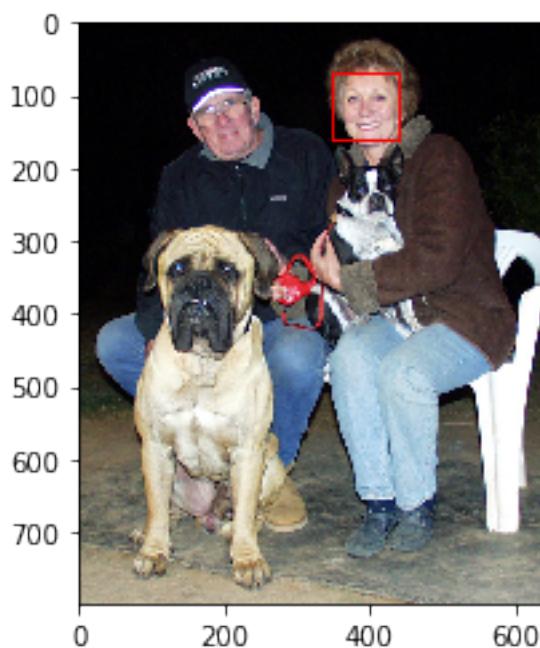
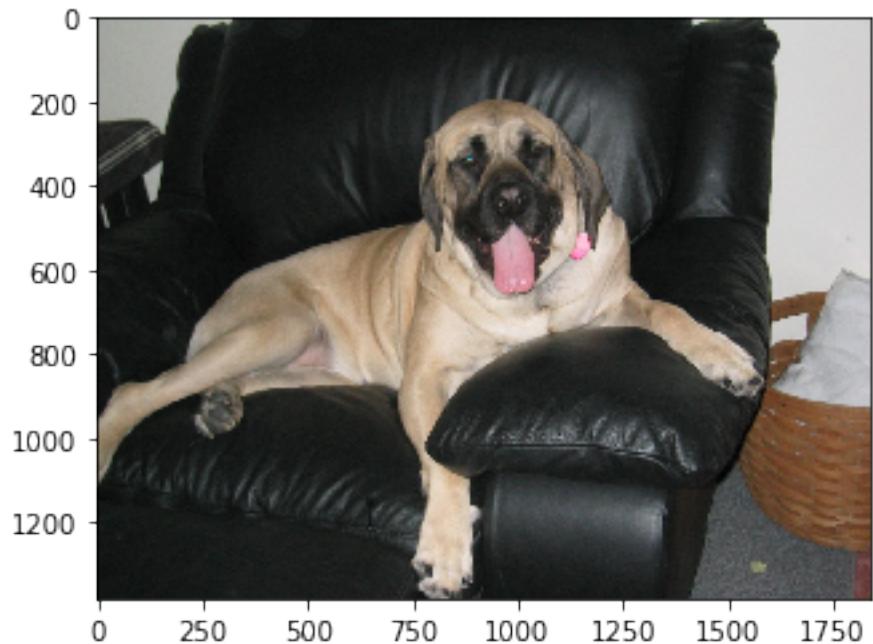
We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

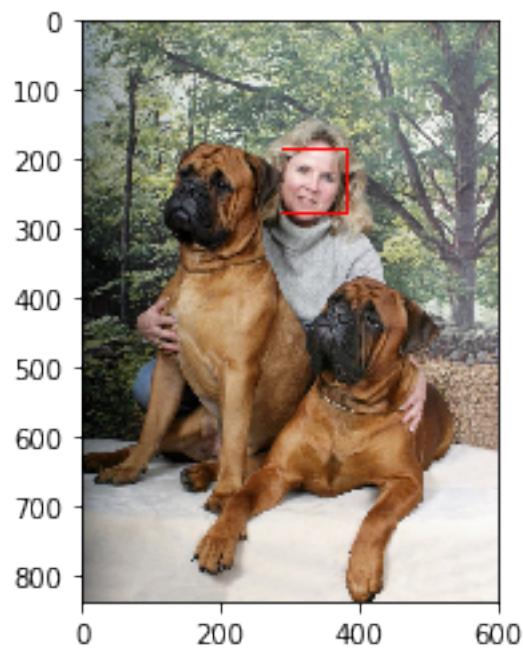
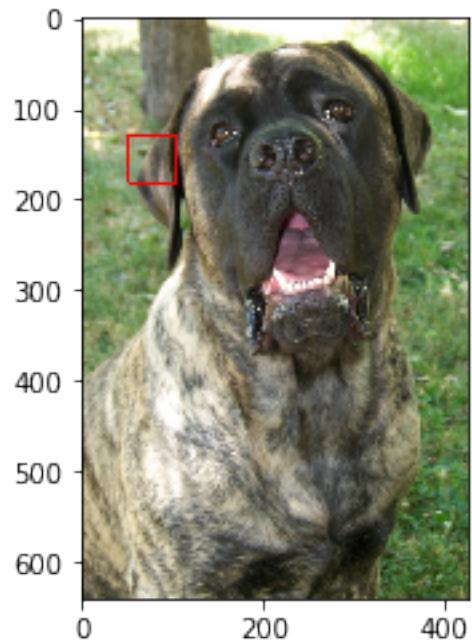
```
In [5]: ### (Optional)
### TODO: Test performance of another face detection algorithm.
### Feel free to use as many code cells as needed.
def face_detector_show(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    # convert BGR image to RGB for plotting
    cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    # display the image, along with bounding box
    for (x,y,w,h) in faces:
        # add bounding box to color image
        cv2.rectangle(cv_rgb,(x,y),(x+w,y+h),(255,0,0),2)
    return len(faces) > 0, cv_rgb
```

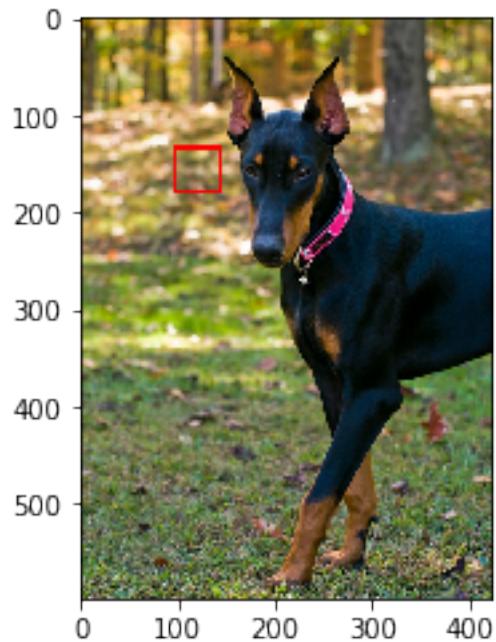
```
In [6]: # Lets check dog images with a detected human face
for i in range(len(human_files_short)):
    detected, img = face_detector_show(dog_files_short[i])
    if detected:
        plt.imshow(img)
        plt.show()
```

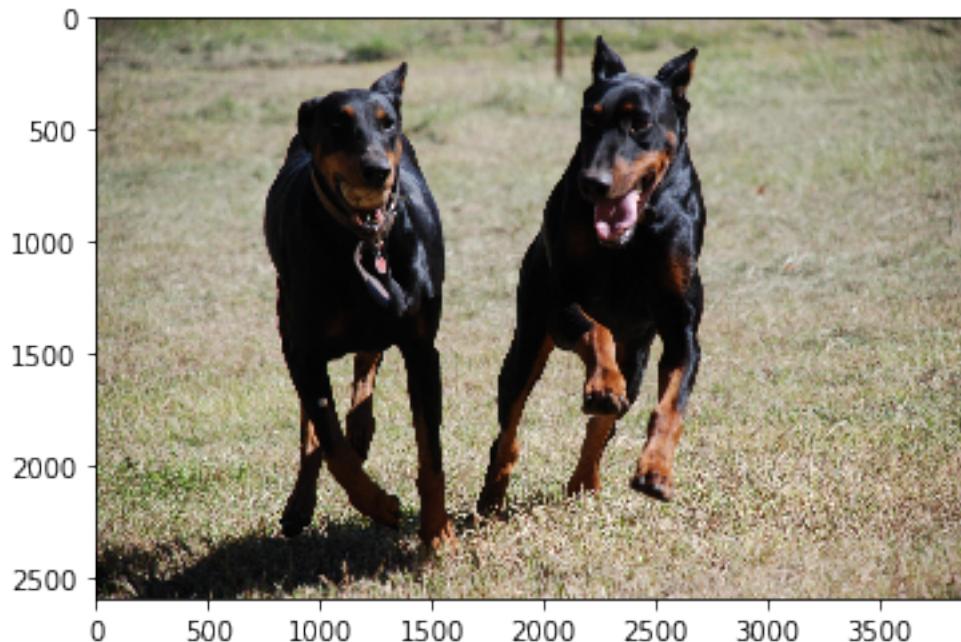




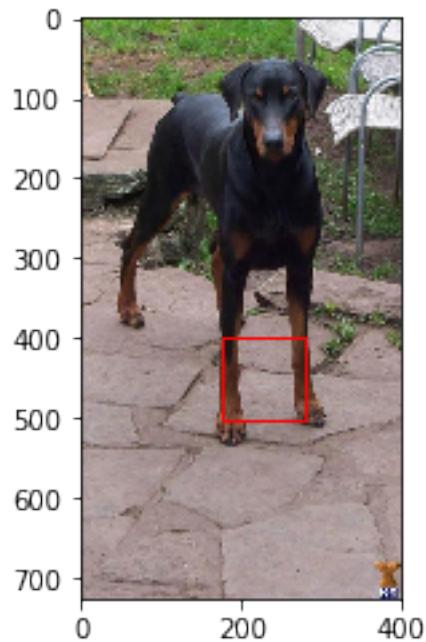


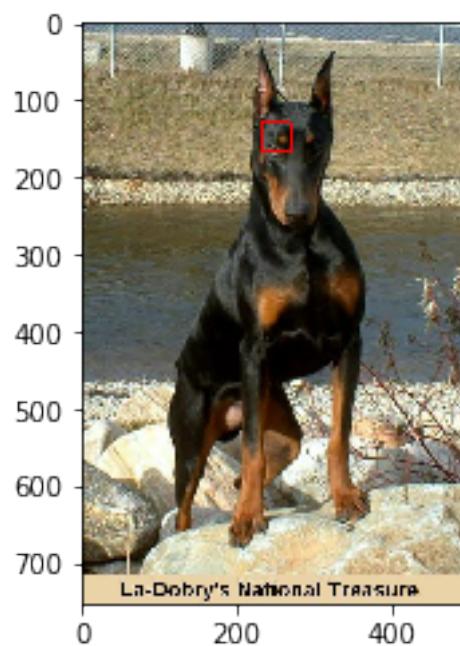
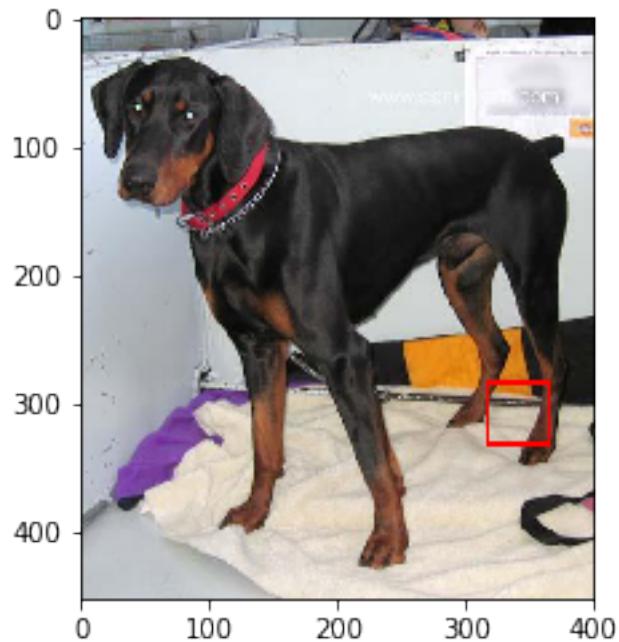






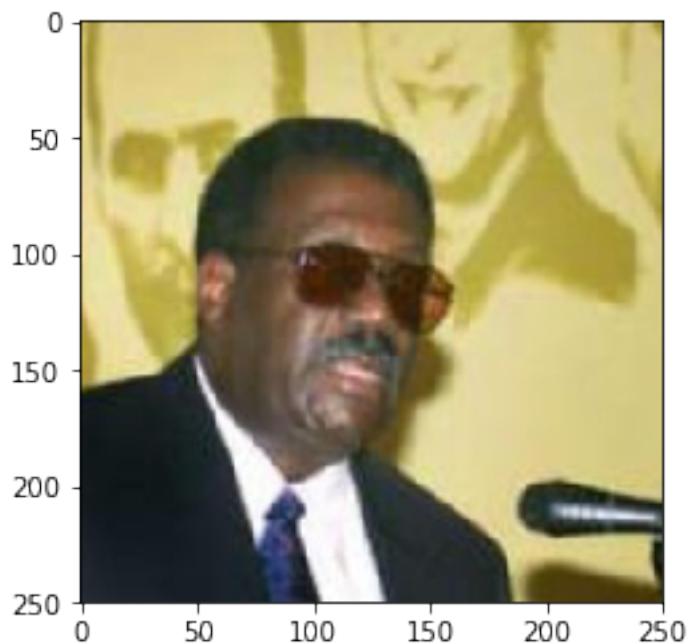
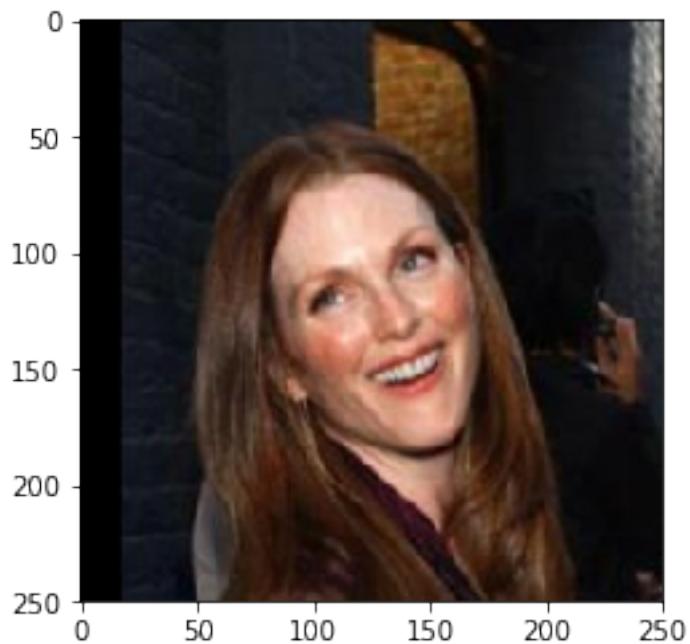






```
In [7]: # Lets check human images which are not detected as human face
for i in range(len(human_files_short)):
    detected, img = face_detector_show(human_files_short[i])
```

```
if not detected:  
    plt.imshow(img)  
    plt.show()
```



---

## ## Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

### 1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](#).

```
In [8]: import torch
        import torchvision.models as models

        # define VGG16 model
        VGG16 = models.vgg16(pretrained=True)

        # check if CUDA is available
        use_cuda = torch.cuda.is_available()

        # move model to GPU if CUDA is available
        if use_cuda:
            VGG16 = VGG16.cuda()
```

```
Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /root/.torch/models/vgg16-397923af.pth
100%|| 553433881/553433881 [00:05<00:00, 94189299.16it/s]
```

```
In [9]: VGG16
```

```
Out[9]: VGG(
    (features): Sequential(
        (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): ReLU(inplace)
        (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (3): ReLU(inplace)
        (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (6): ReLU(inplace)
        (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (8): ReLU(inplace)
        (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (11): ReLU(inplace)
        (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (13): ReLU(inplace)
        (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (15): ReLU(inplace)
```

```

(16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(18): ReLU(inplace)
(19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(20): ReLU(inplace)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(22): ReLU(inplace)
(23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(25): ReLU(inplace)
(26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(27): ReLU(inplace)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace)
(30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(classifier): Sequential(
(0): Linear(in_features=25088, out_features=4096, bias=True)
(1): ReLU(inplace)
(2): Dropout(p=0.5)
(3): Linear(in_features=4096, out_features=4096, bias=True)
(4): ReLU(inplace)
(5): Dropout(p=0.5)
(6): Linear(in_features=4096, out_features=1000, bias=True)
)
)
)

```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

#### 1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as 'dogImages/train/001.Affenpinscher/Affenpinscher\_00001.jpg') as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```

In [10]: from PIL import Image
        from PIL import ImageFile
        ImageFile.LOAD_TRUNCATED_IMAGES = True
        import torchvision.transforms as transforms

def VGG16_predict(img_path):
    """
    Use pre-trained VGG-16 model to obtain index corresponding to
    predicted ImageNet class for image at specified path
    """

```

```

Args:
    img_path: path to an image

Returns:
    Index corresponding to VGG-16 model's prediction
    ...

## TODO: Complete the function.
## Load and pre-process an image from the given img_path
## Return the *index* of the predicted class for that image
image = Image.open(img_path).convert('RGB')

in_transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5),
                      (0.5, 0.5, 0.5))])

img_processed = in_transform(image).unsqueeze(0)

if use_cuda:
    img_processed = img_processed.to('cuda')

VGG16.eval()
output = VGG16(img_processed)
_, top_class = torch.max(output, 1)

return top_class # predicted class index

print(VGG16_predict(dog_files_short[0]))


tensor([ 243], device='cuda:0')

```

### 1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the `dog_detector` function below, which returns `True` if a dog is detected in an image (and `False` if not).

```
In [11]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
```

```

## TODO: Complete the function.
returned_class = VGG16_predict(img_path)
if 151 <= returned_class <= 268:
    return True
else:
    return False      # true/false

```

### 1.1.6 (IMPLEMENTATION) Assess the Dog Detector

**Question 2:** Use the code cell below to test the performance of your dog\_detector function.

- What percentage of the images in `human_files_short` have a detected dog?
- What percentage of the images in `dog_files_short` have a detected dog?

**Answer:** - Percentage of dog images from human files is: 0 - Percentage of dog images from dog files is: 100

```

In [12]: ### TODO: Test the performance of the dog_detector function
          ### on the images in human_files_short and dog_files_short.
detected_dog_from_human_files = 0
detected_dog_from_dog_files = 0
for i in range(len(human_files_short)):
    if dog_detector(human_files_short[i]):
        detected_dog_from_human_files +=1
    if dog_detector(dog_files_short[i]):
        detected_dog_from_dog_files +=1

print("Percentage of dog images from human files is:", detected_dog_from_human_files)
print("Percentage of dog images from dog files is:", detected_dog_from_dog_files)

Percentage of dog images from human files is: 0
Percentage of dog images from dog files is: 100

```

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as Inception-v3, ResNet-50, etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

### ## Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

---

Brittany	Welsh Springer Spaniel
----------	------------------------

---

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

---

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

---

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

---

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

---

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

### 1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms!](#)

In [13]: `!ls /data/dog_images`

`test train valid`

In [14]: `!pwd`

`/home/workspace/dog_project`

In [15]: `import os`  
        `from torchvision import datasets`  
  
        `### TODO: Write data loaders for training, validation, and test sets`  
        `## Specify appropriate transforms, and batch_sizes`  
        `dir_test = os.path.join("/data/dog_images", "test")`

```

dir_valid = os.path.join("/data/dog_images", "valid")
dir_train = os.path.join("/data/dog_images", "train")
train_transforms = transforms.Compose([transforms.RandomRotation(30),
                                      transforms.Resize(225),
                                      transforms.CenterCrop(224),
                                      transforms.RandomHorizontalFlip(),
                                      transforms.ToTensor(),
                                      transforms.Normalize((0.5,), (0.5,))])

valid_transforms = transforms.Compose([transforms.Resize(255),
                                      transforms.CenterCrop(224),
                                      transforms.ToTensor(),
                                      transforms.Normalize((0.5,), (0.5,))])

train_data = datasets.ImageFolder(dir_train, transform=train_transforms)
valid_data = datasets.ImageFolder(dir_valid, transform=valid_transforms)
test_data = datasets.ImageFolder(dir_test, transform=valid_transforms)

train_loader = torch.utils.data.DataLoader(train_data, batch_size=40, shuffle=True)
valid_loader = torch.utils.data.DataLoader(valid_data, batch_size=40)
test_loader = torch.utils.data.DataLoader(test_data, batch_size=40)

loaders_scratch = {
    'train': train_loader,
    'test': test_loader,
    'valid': valid_loader
}

```

In [16]: train\_data.classes

Out[16]: ['001.Affenpinscher',  
          '002.Afghan\_hound',  
          '003.Airedale\_terrier',  
          '004.Akita',  
          '005.Alaskan\_malamute',  
          '006.American\_eskimo\_dog',  
          '007.American\_foxhound',  
          '008.American\_staffordshire\_terrier',  
          '009.American\_water\_spaniel',  
          '010.Anatolian\_shepherd\_dog',  
          '011.Australian\_cattle\_dog',  
          '012.Australian\_shepherd',  
          '013.Australian\_terrier',  
          '014.Basenji',  
          '015.Basset\_hound',  
          '016.Beagle',  
          '017.Bearded\_collie',  
          '018.Beauceron',  
          '019.Bedlington\_terrier',

'020.Belgian\_malinois',  
'021.Belgian\_sheepdog',  
'022.Belgian\_tervuren',  
'023.Bernese\_mountain\_dog',  
'024.Bichon\_frise',  
'025.Black\_and\_tan\_coonhound',  
'026.Black\_russian\_terrier',  
'027.Bloodhound',  
'028.Bluetick\_coonhound',  
'029.Border\_collie',  
'030.Border\_terrier',  
'031.Borzoi',  
'032.Boston\_terrier',  
'033.Bouvier\_des\_flandres',  
'034.Boxer',  
'035.Boykin\_spaniel',  
'036.Briard',  
'037.Brittany',  
'038.Brussels\_griffon',  
'039.Bull\_terrier',  
'040.Bulldog',  
'041.Bullmastiff',  
'042.Cairn\_terrier',  
'043.Canaan\_dog',  
'044.Cane\_corso',  
'045.Cardigan\_welsh\_corgi',  
'046.Cavalier\_king\_charles\_spaniel',  
'047.Chesapeake\_bay\_retriever',  
'048.Chihuahua',  
'049.Chinese\_crested',  
'050.Chinese\_shar-pei',  
'051.Chow\_chow',  
'052.Clumber\_spaniel',  
'053.Cocker\_spaniel',  
'054.Collie',  
'055.Curly-coated\_retriever',  
'056.Dachshund',  
'057.Dalmatian',  
'058.Dandie\_dimmont\_terrier',  
'059.Doberman\_pinscher',  
'060.Dogue\_de\_bordeaux',  
'061.English\_cocker\_spaniel',  
'062.English\_setter',  
'063.English\_springer\_spaniel',  
'064.English\_toy\_spaniel',  
'065.Entlebucher\_mountain\_dog',  
'066.Field\_spaniel',  
'067.Finnish\_spitz',

'068.Flat-coated\_retriever',  
'069.French\_bulldog',  
'070.German\_pinscher',  
'071.German\_shepherd\_dog',  
'072.German\_shorthaired\_pointer',  
'073.German\_wirehaired\_pointer',  
'074.Giant\_schnauzer',  
'075.Glen\_of\_imaal\_terrier',  
'076.Golden\_retriever',  
'077.Gordon\_setter',  
'078.Great\_dane',  
'079.Great\_pyrenees',  
'080.Greater\_swiss\_mountain\_dog',  
'081.Greyhound',  
'082.Havanese',  
'083.Ibizan\_hound',  
'084.Icelandic\_sheepdog',  
'085.Irish\_red\_and\_white\_setter',  
'086.Irish\_setter',  
'087.Irish\_terrier',  
'088.Irish\_water\_spaniel',  
'089.Irish\_wolfhound',  
'090.Italian\_greyhound',  
'091.Japanese\_chin',  
'092.Keeshond',  
'093.Kerry\_blue\_terrier',  
'094.Komondor',  
'095.Kuvasz',  
'096.Labrador\_retriever',  
'097.Lakeland\_terrier',  
'098.Leonberger',  
'099.Lhasa\_apso',  
'100.Lowchen',  
'101.Maltese',  
'102.Manchester\_terrier',  
'103.Mastiff',  
'104.Miniature\_schnauzer',  
'105.Neapolitan\_mastiff',  
'106.Newfoundland',  
'107.Norfolk\_terrier',  
'108.Norwegian\_buhund',  
'109.Norwegian\_elkhound',  
'110.Norwegian\_lundehund',  
'111.Norwich\_terrier',  
'112.Nova\_scotia\_duck\_tolling\_retriever',  
'113.Old\_english\_sheepdog',  
'114.Otterhound',  
'115.Papillon',

```

'116.Parson_russell_terrier',
'117.Pekingese',
'118.Pembroke_welsh_corgi',
'119.Petit_basset_griffon_vendeen',
'120.Pharaoh_hound',
'121.Plott',
'122.Pointer',
'123.Pomeranian',
'124.Poodle',
'125.Portuguese_water_dog',
'126.Saint_bernard',
'127.Silky_terrier',
'128.Smooth_fox_terrier',
'129.Tibetan_mastiff',
'130.Welsh_springer_spaniel',
'131.Wirehaired_pointing_griffon',
'132.Xoloitzcuintli',
'133.Yorkshire_terrier']

```

**Question 3:** Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

**Answer:** First a resize of 225x225 is applied to normalize the training set image, where most of the images will be shrunked, and some few stretched. After a random crop is applied with a size of 224x224. A size of 224 was chosen, similar to what VGG authors use, to try to keep the most information detail possible in the images, while going through multiple pooling layers to reduce its size.

Random flip and rotation is applied to enhance the training process. No translation was applied to avoid decentering the pictures from dog faces

### 1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```

In [17]: import torch.nn as nn
        import torch.nn.functional as F

        # define the CNN architecture
        class Net(nn.Module):
            ### TODO: choose an architecture, and complete the class
            def __init__(self):
                super(Net, self).__init__()
                ## Define layers of a CNN
                hidden_1 = 1024
                hidden_2 = 512
                hidden_3 = 64

```

```

    self.conv11 = nn.Conv2d(3, 16, 3, padding=1)
    self.conv12 = nn.Conv2d(16, 16, 3, padding=1)
    self.conv13 = nn.Conv2d(16, 32, 3, padding=1)
    self.conv2 = nn.Conv2d(32, 32, 3, padding=1)
    self.conv3 = nn.Conv2d(32, 64, 3, padding=1)
    self.conv4 = nn.Conv2d(64, 128, 3, padding=1)
    self.conv5 = nn.Conv2d(128, 256, 3, padding=1)
    self.maxpool = nn.MaxPool2d(2,2)

    self.bn_1 = nn.BatchNorm2d(32)
    self.bn_2 = nn.BatchNorm2d(32)
    self.bn_3 = nn.BatchNorm2d(64)
    self.bn_4 = nn.BatchNorm2d(128)
    self.bn_5 = nn.BatchNorm2d(256)

    self.fc_1 = nn.Linear(256*7*7, hidden_1)
    self.fc_2 = nn.Linear(hidden_1, hidden_2)
    self.fc_3 = nn.Linear(hidden_2, hidden_3)
    self.fc_4 = nn.Linear(hidden_3, 133)

    self.dropout12 = nn.Dropout2d(p = 0.5)
    self.bn1_1 = nn.BatchNorm1d(hidden_1)
    self.bn1_2 = nn.BatchNorm1d(hidden_2)
    self.bn1_3 = nn.BatchNorm1d(hidden_3)

def forward(self, x):
    ## Define forward behavior
    x = F.relu(self.conv11(x))
    x = F.relu(self.conv12(x))
    x = self.bn_1(self.maxpool(F.relu(self.conv13(x))))
    x = self.dropout12(self.bn_2(self.maxpool(F.relu(self.conv2(x)))))

    x = self.bn_3(self.maxpool(F.relu(self.conv3(x))))
    x = self.bn_4(self.maxpool(F.relu(self.conv4(x))))
    x = self.dropout12(self.bn_5(self.maxpool(F.relu(self.conv5(x)))))

    x = x.view(-1, 256*7*7)
    x = self.bn1_1(F.relu(self.fc_1(x)))
    x = self.bn1_2(F.relu(self.fc_2(x)))
    x = self.bn1_3(F.relu(self.fc_3(x)))
    x = F.log_softmax(self.fc_4(x), dim=1)
    return x

def forward(self, x):
    ## Define forward behavior
    x = F.relu(self.conv11(x))
    x = F.relu(self.conv12(x))
    x = self.bn_1(self.maxpool(F.relu(self.conv13(x))))

```

```

        x = self.dropout12(self.bn_2(self.maxpool(F.relu(self.conv2(x)))))
        x = self.bn_3(self.maxpool(F.relu(self.conv3(x))))
        x = self.bn_4(self.maxpool(F.relu(self.conv4(x))))
        x = self.dropout12(self.bn_5(self.maxpool(F.relu(self.conv5(x)))))

        x = x.view(-1, 256*7*7)
        x = self.bn1_1(F.relu(self.fc_1(x)))
        x = self.bn1_2(F.relu(self.fc_2(x)))
        x = self.bn1_3(F.relu(self.fc_3(x)))
        x = F.log_softmax(self.fc_4(x), dim=1)
    return x

```

*#-#-# You so NOT have to modify the code below this line. #-#-#*

```

# instantiate the CNN
model_scratch = Net()

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()

```

In [18]: model\_scratch

Out[18]: Net(

```

(conv11): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(conv12): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(conv13): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(conv2): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(conv3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(conv4): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(conv5): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(maxpool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(bn_1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(bn_2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(bn_3): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(bn_4): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(bn_5): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(fc_1): Linear(in_features=12544, out_features=1024, bias=True)
(fc_2): Linear(in_features=1024, out_features=512, bias=True)
(fc_3): Linear(in_features=512, out_features=64, bias=True)
(fc_4): Linear(in_features=64, out_features=133, bias=True)
(dropout12): Dropout2d(p=0.5)
(bn1_1): BatchNorm1d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(bn1_2): BatchNorm1d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(bn1_3): BatchNorm1d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
```

**Question 4:** Outline the steps you took to get to your final CNN architecture and your reason-

ing at each step.

**Answer:** - The first convolutional layer is defined to take 224 x 224 image having 3 input channels as input and output 16 activation map and then Passed through a relu activation function. The second convolution layer is defined to take 16 activation maps as input and output 16 activation map and then Passed through a relu activation. For the first two convolution layers, Maxpooling layer was not used to retain the image spartial dimensions. - After that convolution layers are defined such a way that the dimension of the image is decreased and the activation map is increased and then convolution layers are passed through the batch normalisation to eliminate randomness and dropout is applied in some of the layers to prevent overfitting. - The output of last convolution layers are flattened and then passed throgh three fully connected layers and passed throgh relu activation function and batch normalisation. - The output of above mentioned fully connected layers are passed through fully connected layer having Log\_softmax activation function for classification.

### 1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

In [19]: `import torch.optim as optim`

```
### TODO: select loss function
criterion_scratch = nn.NLLLoss()

### TODO: select optimizer
optimizer_scratch = optim.Adam(model_scratch.parameters(), lr=0.001)
```

### 1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath '`model_scratch.pt`'.

In [20]: `def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):`  
    `"""returns trained model"""`  
    `# initialize tracker for minimum validation loss`  
    `valid_loss_min = np.Inf`  
  
    `for epoch in range(1, n_epochs+1):`  
        `# initialize variables to monitor training and validation loss`  
        `train_loss = 0.0`  
        `valid_loss = 0.0`  
  
        `#####`  
        `# train the model #`  
        `#####`  
        `model.train()`  
        `for batch_idx, (data, target) in enumerate(loaders['train']):`  
            `# move to GPU`  
            `if use_cuda:`

```

        data, target = data.cuda(), target.cuda()
        optimizer.zero_grad()
        logp = model.forward(data)
        loss = criterion(logp, target)
        loss.backward()
        optimizer.step()
        ## find the loss and update the model parameters accordingly
        ## record the average training loss, using something like
        train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

#####
# validate the model #
#####
model.eval()
for batch_idx, (data, target) in enumerate(loaders['valid']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    logp = model.forward(data)
    loss = criterion(logp, target)
    valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (loss.data - valid_loss))
    ## update the average validation loss

# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
    epoch,
    train_loss,
    valid_loss
))

## TODO: save the model if validation loss has decreased
if valid_loss <= valid_loss_min:
    torch.save(model.state_dict(), save_path)
    valid_loss_min = valid_loss
    print('.....initiating save mode')
# return trained model
return model

# train the model
model_scratch = train(25, loaders_scratch, model_scratch, optimizer_scratch,
                      criterion_scratch, use_cuda, 'model_scratch.pt')

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('model_scratch.pt'))

```

Epoch: 1            Training Loss: 4.807509            Validation Loss: 4.687277

```

...initiating save mode
Epoch: 2          Training Loss: 4.471549          Validation Loss: 4.404094
...initiating save mode
Epoch: 3          Training Loss: 4.311878          Validation Loss: 4.426095
Epoch: 4          Training Loss: 4.193388          Validation Loss: 4.248815
...initiating save mode
Epoch: 5          Training Loss: 4.075240          Validation Loss: 4.127416
...initiating save mode
Epoch: 6          Training Loss: 3.998126          Validation Loss: 4.019804
...initiating save mode
Epoch: 7          Training Loss: 3.889716          Validation Loss: 4.318954
Epoch: 8          Training Loss: 3.815179          Validation Loss: 3.881942
...initiating save mode
Epoch: 9          Training Loss: 3.740013          Validation Loss: 3.858064
...initiating save mode
Epoch: 10         Training Loss: 3.647617          Validation Loss: 3.897732
Epoch: 11         Training Loss: 3.564791          Validation Loss: 3.701886
...initiating save mode
Epoch: 12         Training Loss: 3.450893          Validation Loss: 3.886198
Epoch: 13         Training Loss: 3.383608          Validation Loss: 3.572875
...initiating save mode
Epoch: 14         Training Loss: 3.253147          Validation Loss: 3.596771
Epoch: 15         Training Loss: 3.168374          Validation Loss: 3.429441
...initiating save mode
Epoch: 16         Training Loss: 3.052519          Validation Loss: 3.562726
Epoch: 17         Training Loss: 2.943792          Validation Loss: 3.418638
...initiating save mode
Epoch: 18         Training Loss: 2.837054          Validation Loss: 3.274237
...initiating save mode
Epoch: 19         Training Loss: 2.744966          Validation Loss: 3.198236
...initiating save mode
Epoch: 20         Training Loss: 2.621045          Validation Loss: 3.225266
Epoch: 21         Training Loss: 2.522975          Validation Loss: 3.151161
...initiating save mode
Epoch: 22         Training Loss: 2.412591          Validation Loss: 3.069539
...initiating save mode
Epoch: 23         Training Loss: 2.310749          Validation Loss: 2.995447
...initiating save mode
Epoch: 24         Training Loss: 2.181733          Validation Loss: 3.067174
Epoch: 25         Training Loss: 2.071960          Validation Loss: 3.049443

```

### 1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```
In [21]: def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['test']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
        100. * correct / total, correct, total))

    # call test function
    test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)
```

Test Loss: 2.984463

Test Accuracy: 27% (231/836)

## Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)  
 You will now use transfer learning to create a CNN that can identify dog breed from images.  
 Your CNN must attain at least 60% accuracy on the test set.

### 1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate **data loaders** for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

In [22]: *## TODO: Specify data loaders*

```
dir_test = os.path.join("/data/dog_images", "test")
dir_valid = os.path.join("/data/dog_images", "valid")
dir_train = os.path.join("/data/dog_images", "train")
train_transforms_dog = transforms.Compose([transforms.RandomRotation(30),
                                         transforms.Resize(225),
                                         transforms.CenterCrop(224),
                                         transforms.RandomHorizontalFlip(),
                                         transforms.ToTensor(),
                                         transforms.Normalize((0.5,), (0.5,))])

valid_transforms_dog = transforms.Compose([transforms.Resize(255),
                                         transforms.CenterCrop(224),
                                         transforms.ToTensor(),
                                         transforms.Normalize((0.5,), (0.5,))])
train_data_dog = datasets.ImageFolder(dir_train, transform=train_transforms_dog)
valid_data_dog = datasets.ImageFolder(dir_valid, transform=valid_transforms_dog)
test_data_dog = datasets.ImageFolder(dir_test, transform=valid_transforms_dog)

train_loader_dog = torch.utils.data.DataLoader(train_data_dog, batch_size=40, shuffle=True)
valid_loader_dog = torch.utils.data.DataLoader(valid_data_dog, batch_size=40)
test_loader_dog = torch.utils.data.DataLoader(test_data_dog, batch_size=40)

loaders_transfer={'train':train_loader_dog,
                  'valid': valid_loader_dog,
                  'test': test_loader_dog}
```

### 1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

In [23]: `import torchvision.models as models  
import torch.nn as nn`

```
## TODO: Specify model architecture
model_transfer = models.resnet50(pretrained=True)
for params in model_transfer.parameters():
    params.require_grad = True
model_transfer.fc = nn.Linear(2048, 133)

if use_cuda:
    model_transfer = model_transfer.cuda()
```

```
Downloading: "https://download.pytorch.org/models/resnet50-19c8e357.pth" to /root/.torch/models/
100%|| 102502400/102502400 [00:01<00:00, 72648100.20it/s]
```

```
In [24]: model_transfer
```

```
Out[24]: ResNet(
    (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
    (layer1): Sequential(
        (0): Bottleneck(
            (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (relu): ReLU(inplace)
            (downsample): Sequential(
                (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            )
        )
        (1): Bottleneck(
            (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (relu): ReLU(inplace)
        )
        (2): Bottleneck(
            (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (relu): ReLU(inplace)
        )
    )
    (layer2): Sequential(
        (0): Bottleneck(
            (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        )
    )
)
```

```

(conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias
(bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats
(conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats
(relu): ReLU(inplace)
(downsample): Sequential(
    (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats
)
)
(1): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats
    (relu): ReLU(inplace)
)
(2): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats
    (relu): ReLU(inplace)
)
(3): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats
    (relu): ReLU(inplace)
)
)
(layer3): Sequential(
(0): Bottleneck(
    (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats
    (relu): ReLU(inplace)
    (downsample): Sequential(

```

```

        (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
)
(1): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
)
(2): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
)
(3): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
)
(4): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
)
(5): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
)

```

```

        )
    )
(layer4): Sequential(
    (0): Bottleneck(
        (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace)
        (downsample): Sequential(
            (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        )
    )
    (1): Bottleneck(
        (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace)
    )
    (2): Bottleneck(
        (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace)
    )
)
(avgpool): AvgPool2d(kernel_size=7, stride=1, padding=0)
(fc): Linear(in_features=2048, out_features=133, bias=True)
)

```

**Question 5:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

**Answer:** - Resnet50 is used because it is relatively fast to train - Resnet50 pretrained weights are used because it was trained on similar data as the train set - In final classification layers only the last layer is changed to a dense linear with the same number of nodes as the number of classes of dogs which we are trying to predict. That is the part of the network which I am going to re-train so that it is able to classify the dog breeds. - This architecture is suitable because the training data is relatively large and has similar data to what the original pretrained model was trained on. - The weights are randomly initialized in the new fully connected layer and in the rest of the network,

the weights are initialized using the pre-trained weights and re-train the entire neural network

### 1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```
In [25]: criterion_transfer = nn.CrossEntropyLoss()  
optimizer_transfer = optim.SGD(model_transfer.parameters(), lr=0.01)
```

### 1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath '`model_transfer.pt`'.

```
In [26]: # train the model  
n_epochs = 15  
model_transfer = train(n_epochs , loaders_transfer, model_transfer, optimizer_transfer,  
  
# load the model that got the best validation accuracy (uncomment the line below)  
model_transfer.load_state_dict(torch.load('model_transfer.pt'))
```

Epoch: 1	Training Loss: 3.808256	Validation Loss: 2.297208
...initiating save mode		
Epoch: 2	Training Loss: 1.895583	Validation Loss: 1.200905
...initiating save mode		
Epoch: 3	Training Loss: 1.139708	Validation Loss: 0.815288
...initiating save mode		
Epoch: 4	Training Loss: 0.816437	Validation Loss: 0.647269
...initiating save mode		
Epoch: 5	Training Loss: 0.632947	Validation Loss: 0.538353
...initiating save mode		
Epoch: 6	Training Loss: 0.511475	Validation Loss: 0.471149
...initiating save mode		
Epoch: 7	Training Loss: 0.436376	Validation Loss: 0.435332
...initiating save mode		
Epoch: 8	Training Loss: 0.365064	Validation Loss: 0.418808
...initiating save mode		
Epoch: 9	Training Loss: 0.319853	Validation Loss: 0.391794
...initiating save mode		
Epoch: 10	Training Loss: 0.279545	Validation Loss: 0.391162
...initiating save mode		
Epoch: 11	Training Loss: 0.247332	Validation Loss: 0.361073
...initiating save mode		
Epoch: 12	Training Loss: 0.224221	Validation Loss: 0.354447
...initiating save mode		
Epoch: 13	Training Loss: 0.196717	Validation Loss: 0.371353
Epoch: 14	Training Loss: 0.170017	Validation Loss: 0.333706
...initiating save mode		

```
Epoch: 15          Training Loss: 0.153742          Validation Loss: 0.327975
...initiating save mode
```

### 1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [27]: test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

```
Test Loss: 0.337666
```

```
Test Accuracy: 89% (748/836)
```

### 1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
In [28]: data_transfer = loaders_transfer
```

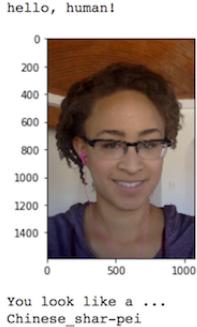
```
In [29]: ### TODO: Write a function that takes a path to an image as input
        ### and returns the dog breed that is predicted by the model.
        # list of class names by index, i.e. a name can be accessed like class_names[0]
        from torch.autograd import Variable
        class_names = [item[4:].replace("_", " ") for item in train_data_dog.classes]

        def predict_breed_transfer(img_path):
            # load the image and return the predicted breeds for that image
            image = Image.open(img_path).convert('RGB')

            transformer = transforms.Compose([transforms.Resize(255),
                                              transforms.CenterCrop(224),
                                              transforms.ToTensor(),
                                              transforms.Normalize((0.5,), (0.5,))])

            img_processed = transformer(image).unsqueeze(0)

            if use_cuda:
                img_processed = img_processed.to('cuda')
            img_var = Variable(img_processed, requires_grad= False)
            img_var = img_var.cuda()
            model_transfer.eval()
            output = model_transfer(img_var)
            _, pred = torch.max(output, 1)
            return class_names[pred]
```



Sample Human Output

```
prediction = predict_breed_transfer(dog_files[1002])
print(prediction)
```

Manchester terrier

#### ## Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

#### 1.1.18 (IMPLEMENTATION) Write your Algorithm

In [30]: *### TODO: Write your algorithm.*  
*### Feel free to use as many code cells as needed.*

```
def run_app(img_path):
    ## handle cases for a human face, dog, and neither
    image = Image.open(img_path).convert('RGB')
    plt.imshow(image)
    plt.show()
    if dog_detector(img_path) == True:
        print('Dog found! Predicted breed is: ' + predict_breed_transfer(img_path))

    elif face_detector(img_path) == True:
        print("Human found! Most similar breed is: " + predict_breed_transfer(img_path))
```

```
    else:  
        print('No human or dog was found')
```

---

## ## Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

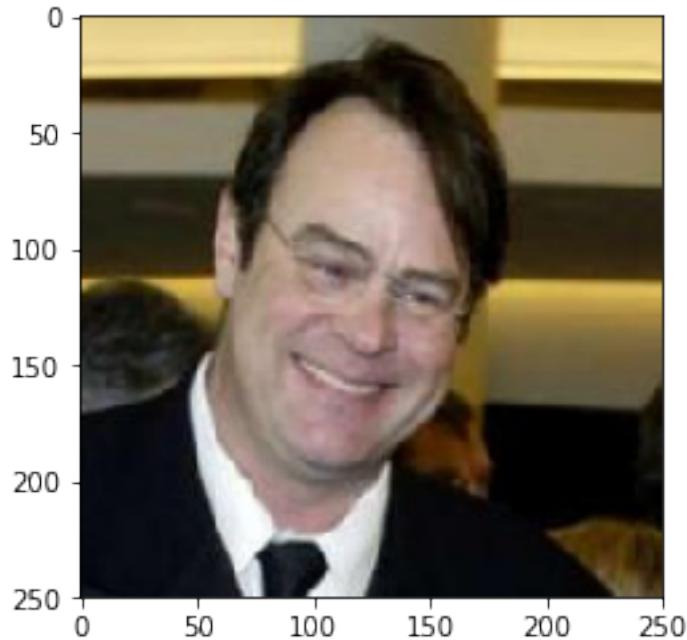
### 1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

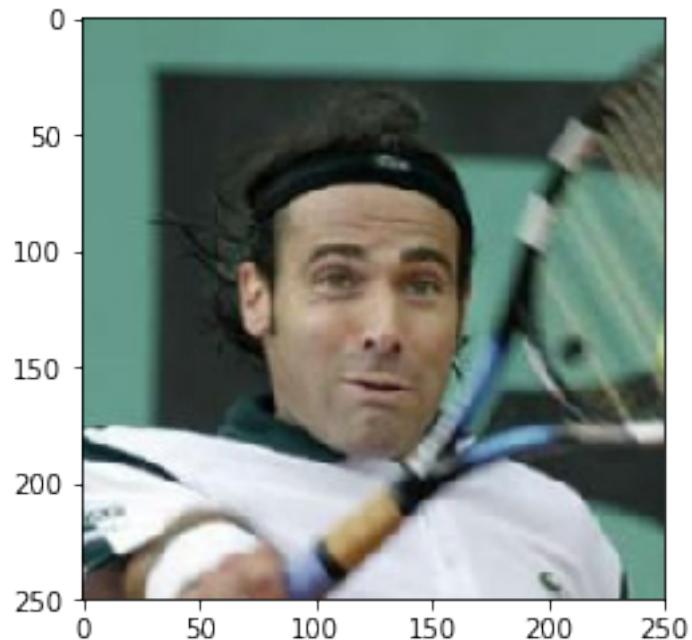
**Question 6:** Is the output better than you expected :) ? Or worse :( ? Provide at least three possible points of improvement for your algorithm.

**Answer:** - The breed detector could be improved by changing the classifier network to increase accuracy. - Some training dataset for some dog breed have very less images which can be added with more images to get better accuracy - Train the predictor to handle images with both human and dog

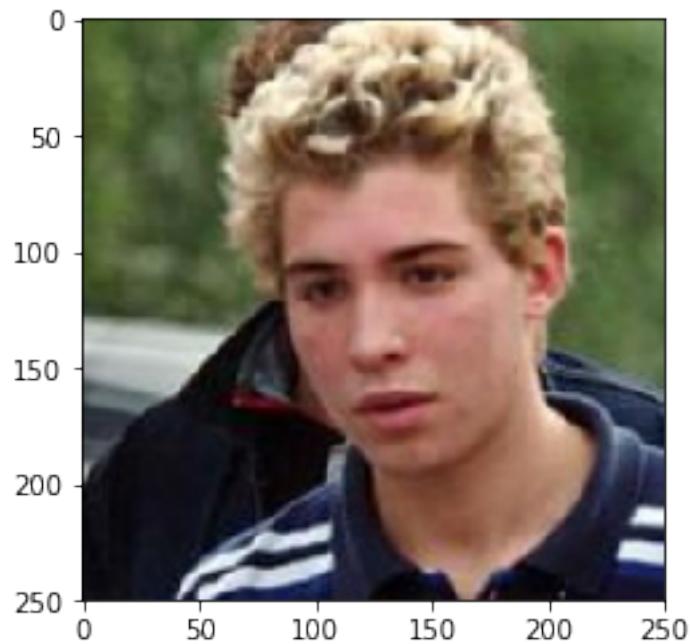
```
In [31]: ## TODO: Execute your algorithm from Step 6 on  
## at least 6 images on your computer.  
## Feel free to use as many code cells as needed.  
  
## suggested code, below  
for file in np.hstack((human_files[:3], dog_files[:3])):  
    run_app(file)
```



Human found! Most similar breed is: Chihuahua



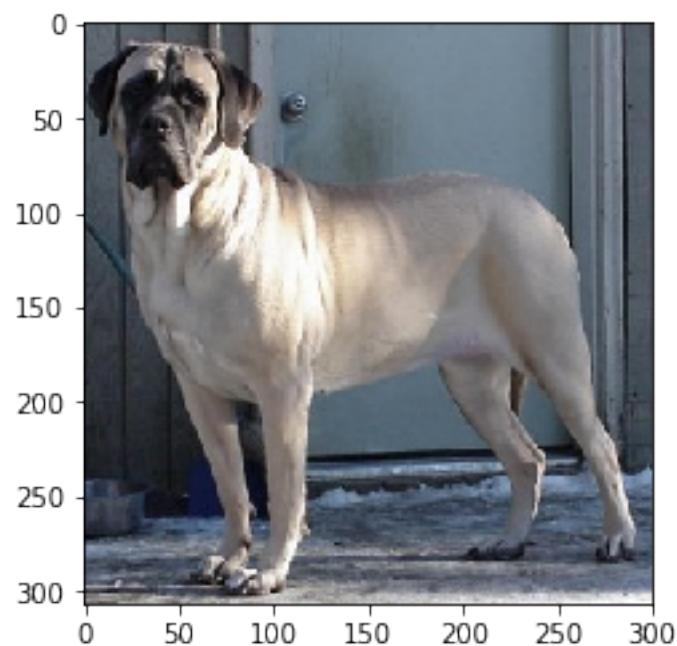
Human found! Most similar breed is: Neapolitan mastiff



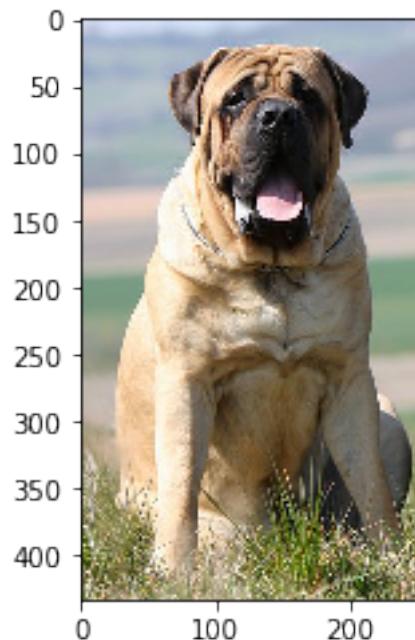
Human found! Most similar breed is: American water spaniel



Dog found! Predicted breed is: Mastiff



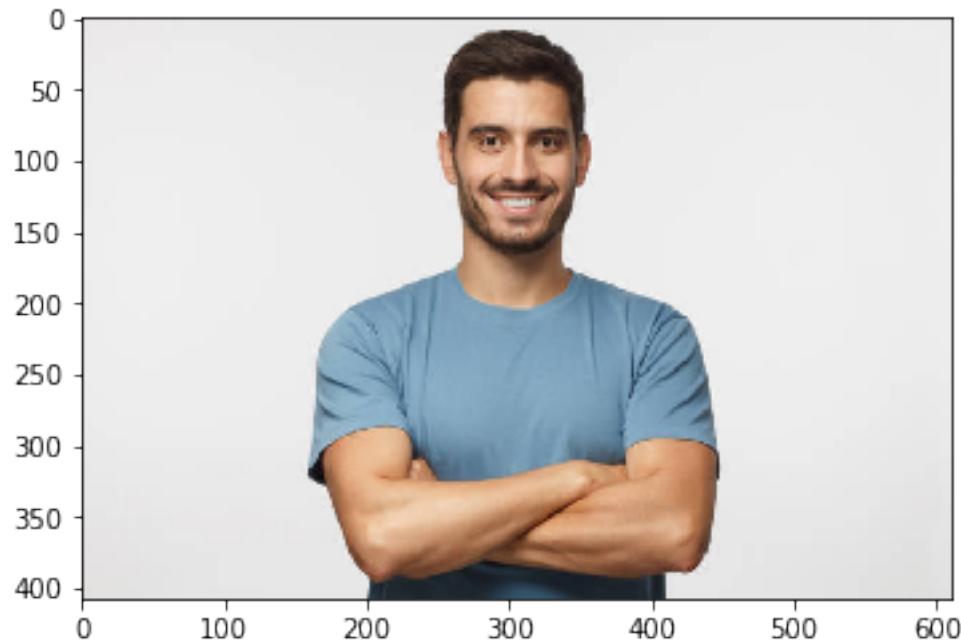
Dog found! Predicted breed is: Mastiff



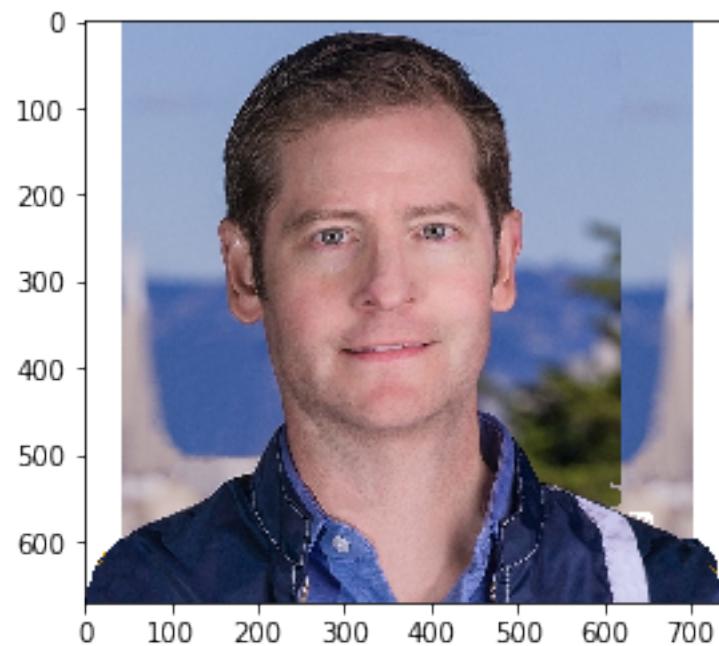
Dog found! Predicted breed is: Mastiff

```
In [37]: my_human_files = ['./my_images/human3.jpg', './my_images/Accenture-Human-Machine-AI-Jam']  
my_dog_files = ['./my_images/American_water_spaniel_00648-Copy1.jpg', './my_images/British']
```

```
In [38]: for file in np.hstack((my_human_files, my_dog_files)):  
    run_app(file)
```



Human found! Most similar breed is: Chinese shar-pei



Human found! Most similar breed is: Pointer



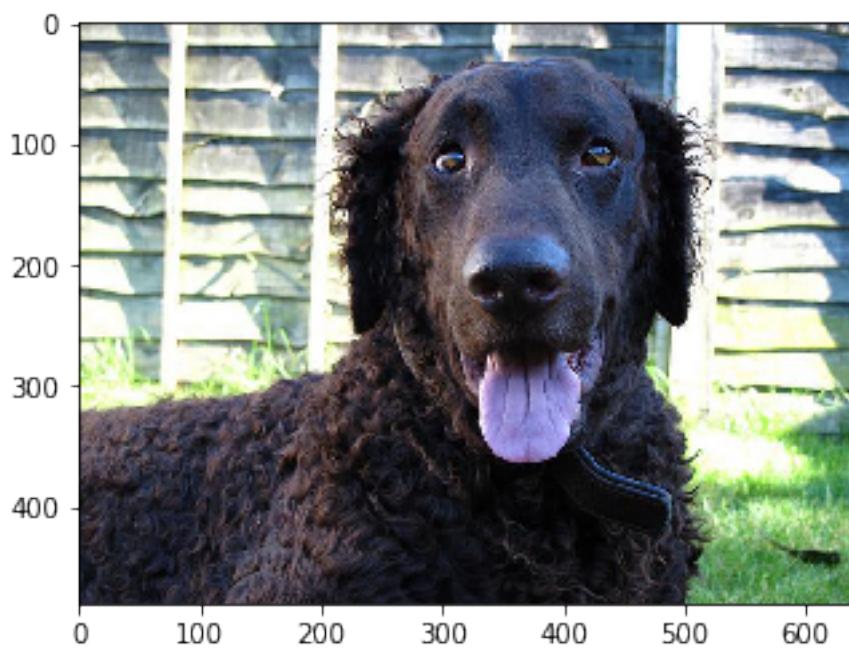
Human found! Most similar breed is: Dogue de bordeaux



Dog found! Predicted breed is: Curly-coated retriever



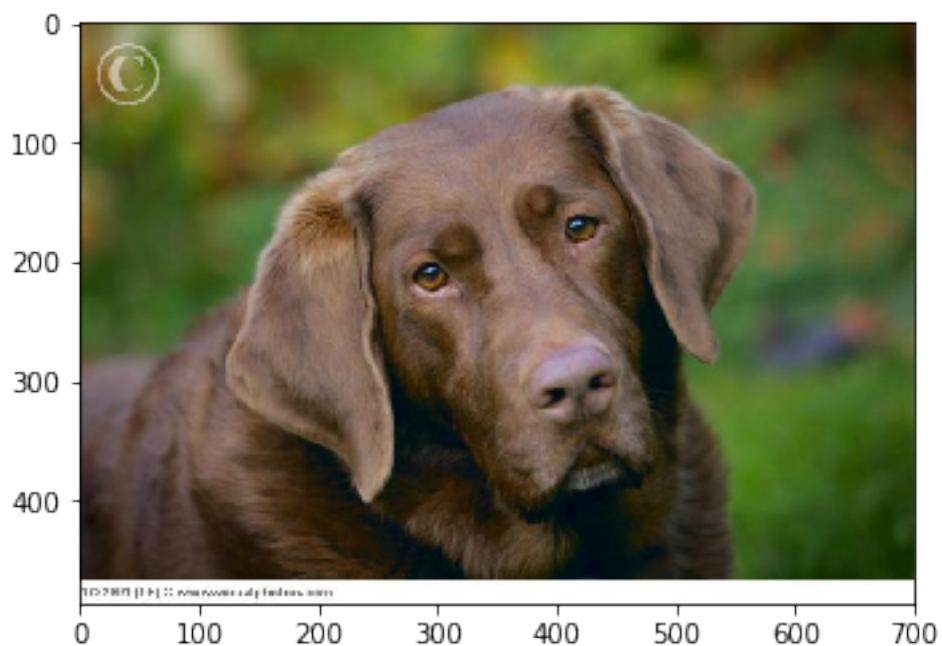
Dog found! Predicted breed is: Brittany



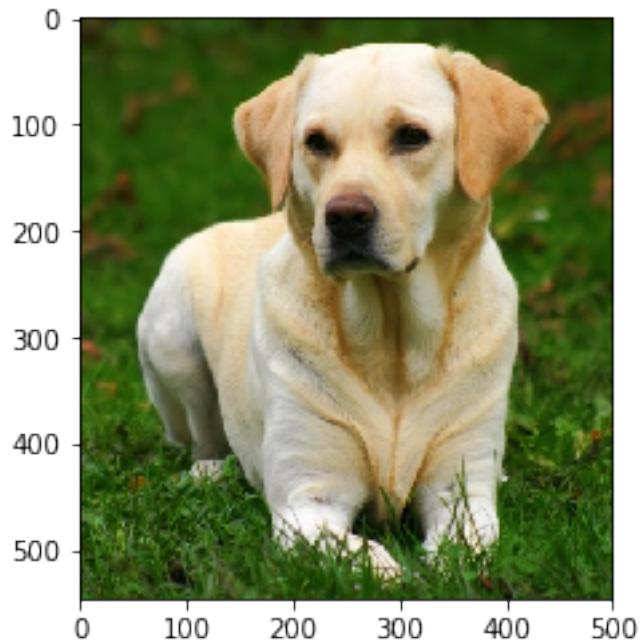
Dog found! Predicted breed is: Curly-coated retriever



Dog found! Predicted breed is: Labrador retriever



Dog found! Predicted breed is: Labrador retriever



Dog found! Predicted breed is: Labrador retriever



Dog found! Predicted breed is: Welsh springer spaniel

In [ ]: